

Information Leakage in Camouflage and Comparison to Our System

Concurrent work by Zhou et al. introduces Camouflage [1], a memory traffic-shaping technique that builds on MITTS [2] to mitigate timing channel attacks through memory controllers. In this section, we discuss fundamental limitations that prevent Camouflage from offering provable bounds on information leakage and demonstrate a covert channel that breaks Camouflage at bandwidths of 10s of Kbps.

Camouflage design. The baseline MITTS design ensures that an application’s memory traffic does not exceed a specified bandwidth distribution *per replenishment time period*. This distribution is a histogram of latency between two consecutive memory requests, with `credits` assigned to each bin of the histogram. When a memory request misses in the per-core cache, the time delta between the current request and the previous request is calculated and is checked against the bandwidth histogram. If the corresponding bin has credits left in the histogram, the credit is deducted from the histogram and the request is issued to the uncore – i.e., MITTS *drains* the histogram until credits remain. At the end of each replenishment time period, the histogram is replenished to its original state.

Camouflage leverages the hardware proposed in MITTS to shape traffic as a means of thwarting side and covert channel attacks. To further obfuscate contention for the underlying resource, Camouflage introduces *fake memory requests* on behalf of workloads that do not fully drain a histogram during a replenishment period. At the end of a replenishment period, the remaining credits are stored and sent out during the subsequent time replenishment period. These fake memory requests are sent at a lower priority than the program’s real memory requests to not interfere with a program’s performance.

Camouflage’s fundamental limitations. However, a significant limitation of the method in Camouflage is ignoring dependencies across time. The random variable X in Camouflage, which corresponds to inter-arrival times obtained from program traces, actually has strong time dependency: memory traces are well-known to have bursts and hence small inter-arrival times likely follow other small inter-arrival times. A sequence of time-dependent random variables is called a *stochastic process*. Therefore, Camouflage’s actual problem consists of estimating entropy and mutual information of two stochastic processes $X(t), Y(t)$. This is a significantly more challenging problem compared to estimating the entropy and mutual information of two *scalar* random variables X, Y from independent samples (which is how Camouflage estimates information leakage) [1, 3, 4, 5].

The problem of estimating the entropy rates of stochastic

processes has been studied extensively (see [3] Chapter 4 for an accessible introduction and [4, 5] and references therein for various algorithms used for different types of stochastic processes). Without making any assumptions on the type of memory in the stochastic process $X(t)$, estimating entropy requires computing histograms of all possible combinations of t values, which scales exponentially in the number of observed timeslots.

Camouflage ignores time dependencies and claims to prevent leakage only over long-term timing information for which, perhaps, time dependencies are not as strong. Unfortunately, the time-scale for which trace time dependencies are not significant depends on the application under test and is very hard to quantify. Worse, such time dependencies cannot be ignored if the application creates a *covert channel*. We use this time dependencies to construct a high bandwidth covert channel.

Covert channel. Camouflage assumes that an adversary or malicious process cannot determine timing information within the *replenishment period* over which it shapes memory traffic [1]. Therefore, we create a covert channel where the communicating parties can see timing information over a time interval *larger* than the replenishment period of Camouflage. We transmit one bit over every such *communication interval* through a memory channel shaped by Camouflage. We additionally assert that malicious processes may learn the enforced distribution by profiling over the communication interval since Camouflage must shape workloads to a fixed inter-arrival time distribution [6]. We assume communication intervals on the order of 1,000 memory cycles [1] and a memory clock speed on the order of hundreds of MHz. Sending 1 bit per communication interval creates a covert channel with bandwidth on the order of 10 Kbps.

While transmission parameters may be optimized further, we assume that the communication interval is equal to two or more replenishment periods. A *transmitter* sends one bit by draining the learned histograms either partially or fully over the communication interval and letting the receiver sense the difference. At the same time, a *receiver* attempts to drain its histograms fully during each communication window. The transmitter draining its histograms fully ensures that the receiver will perceive contention throughout the communication window, which it measures by observing the average latency of its memory requests. If the transmitter only partially drains its histograms, there will be less contention throughout the interval, as Camouflage only introduces fake accesses retroactively in the next replenishment window. As the communication

interval encompasses multiple such replenishment windows, fake requests due to the transmitter partially draining its histograms will not be present during a subsequent communication interval. The receiver may therefore determine a threshold on the latency it perceives to distinguish the bits being sent over the channel.

Shaping in our system. In this paper, we explicitly handle time dependencies and prove precise bounds on the amount of leaked information *even if the adversary has perfectly precise timing information at any resolution*. The key reason (explained in further detail in Sections ?? and ??) is the following: Camouflage is *draining* a histograms of requests and allowing credits from one replenishment window to affect an arbitrary number of following windows – this allows time dependencies from $X(t)$ to transfer to $Y(t)$. Hence, Camouflage fundamentally leaks information and is very hard to test or control for in an application. On the contrary, our approach is to learn a distribution for each time slot t in a program’s execution, and then *sample* $Y(t)$ from these distributions. This allows us to generate a masked trace $Y(t)$ that is independent across time (i.e., is secure), but not identically distributed (and hence improves performance). By learning the statistics of $X(t)$ for each time t we ensure that our samples are quite close to the real trace, but privacy can be controlled by exploiting independence across time for the process $Y(t)$ that the adversary observes.

References

- [1] Yanqi Zhou, Sameer Wagh, Prateek Mittal, and David Wentzlaff. Camouflage: Memory traffic shaping to mitigate timing attacks. In *2016 IEEE International Symposium on High Performance Computer Architecture, HPCA 2017, HPCA ’17*, Piscataway, NJ, USA, 2017. IEEE Press.
- [2] Yanqi Zhou and David Wentzlaff. Mitts: Memory inter-arrival time traffic shaping. In *Proceedings of the 43rd International Symposium on Computer Architecture, ISCA ’16*, pages 532–544, Piscataway, NJ, USA, 2016. IEEE Press.
- [3] Thomas M. Cover and Joy A. Thomas. *Elements of Information Theory (Wiley Series in Telecommunications and Signal Processing)*. Wiley-Interscience, 2006.
- [4] Yun Gao, Ioannis Kontoyiannis, and Elie Bienenstock. Estimating the entropy of binary time series: Methodology, some theory and a simulation study. *Entropy*, 10(2):71–99, 2008.
- [5] Yun Gao, Ioannis Kontoyiannis, and Elie Bienenstock. From the entropy to the statistical structure of spike trains. *CoRR*, abs/0710.4117, 2007.
- [6] C. Hunger, M. Kazdagli, A. Rawat, A. Dimakis, S. Vishwanath, and M. Tiwari. Understanding contention-based channels and using them for defense. In *High Performance Computer Architecture (HPCA), 2015 IEEE 21st International Symposium on*, pages 639–650, Feb 2015.